

ISOVALENT

Better Bandwidth Management with eBPF

Daniel Borkmann

Christopher M. Luciano

Nikolay Aleksandrov



Problem Statement


- Increasing Pod density per node
- Competition for node resources e.g. CPU and memory
- Optimization problem for operators: Resource allocation and efficient use, achieving SLOs, etc

```
apiVersion: v1
kind: Pod
metadata:
  name: frontend
spec:
  containers:
  - name: app
    image: images.my-company.example/app:v4
    resources:
      requests:
        memory: "64Mi"
        cpu: "250m"
      limits:
        memory: "128Mi"
        cpu: "500m"
```


What the container is guaranteed to get, e.g. kubelet will only schedule the Pod on a node which can provide this resource.

Hard upper limit, ensures that container never goes above this threshold.

Pods per Node




Median Containers per Host



Source: Sysdig 2022 Cloud Native Security and Usage Report

Problem Statement

- But what about networking?
- TCP sends AFAP (as fast as possible)
- AFAP output contract, shaping typically implemented by device output queues
- Queue length limit & receive window determines in-flight rate
- “How fast” implicit in queue drain rate
- Who limits a Pod’s network usage in Kubernetes?



Source: <https://hpbn.co/building-blocks-of-tcp/>

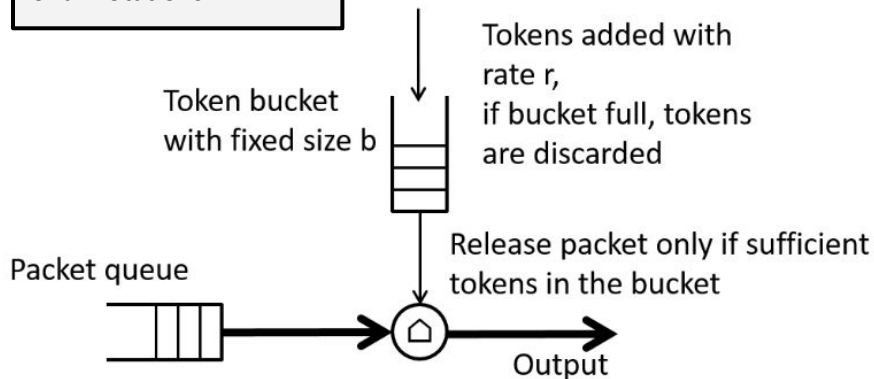


Problem Statement

- Kubernetes bandwidth enforcement has only been experimental so far :- (
- Support for Pod annotations with 'outsourced' bandwidth meta plugin

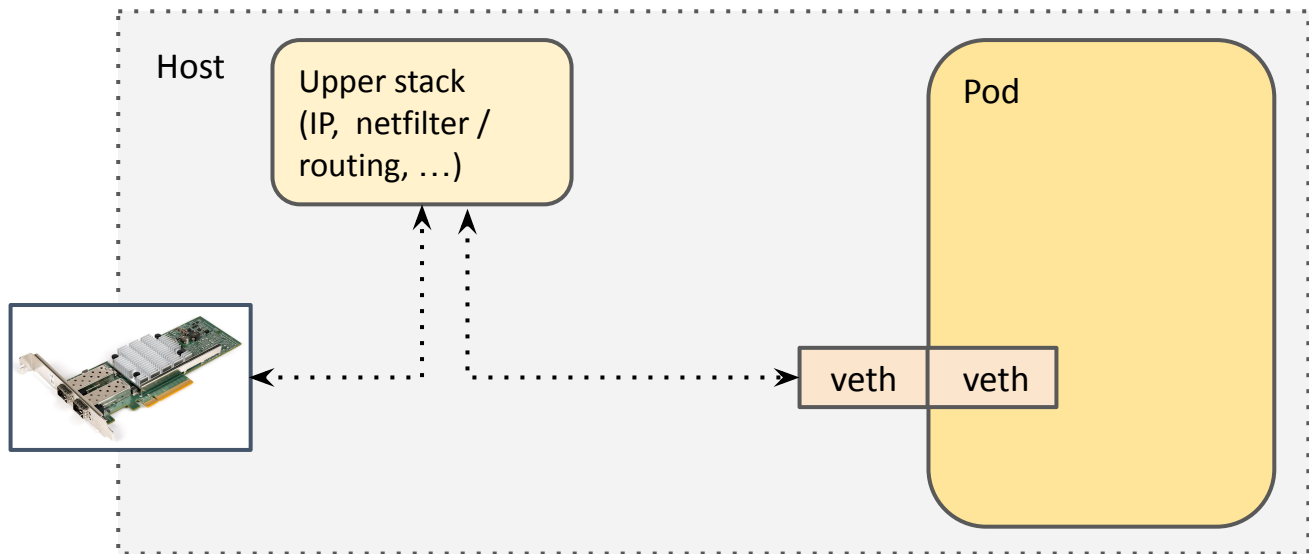
```
apiVersion: v1
kind: Pod
metadata:
  annotations:
    kubernetes.io/ingress-bandwidth: 1M
    kubernetes.io/egress-bandwidth: 1M
...
```

Adds rudimentary **token bucket filters (TBF)** to implement enforcement of annotations.




Problem Statement

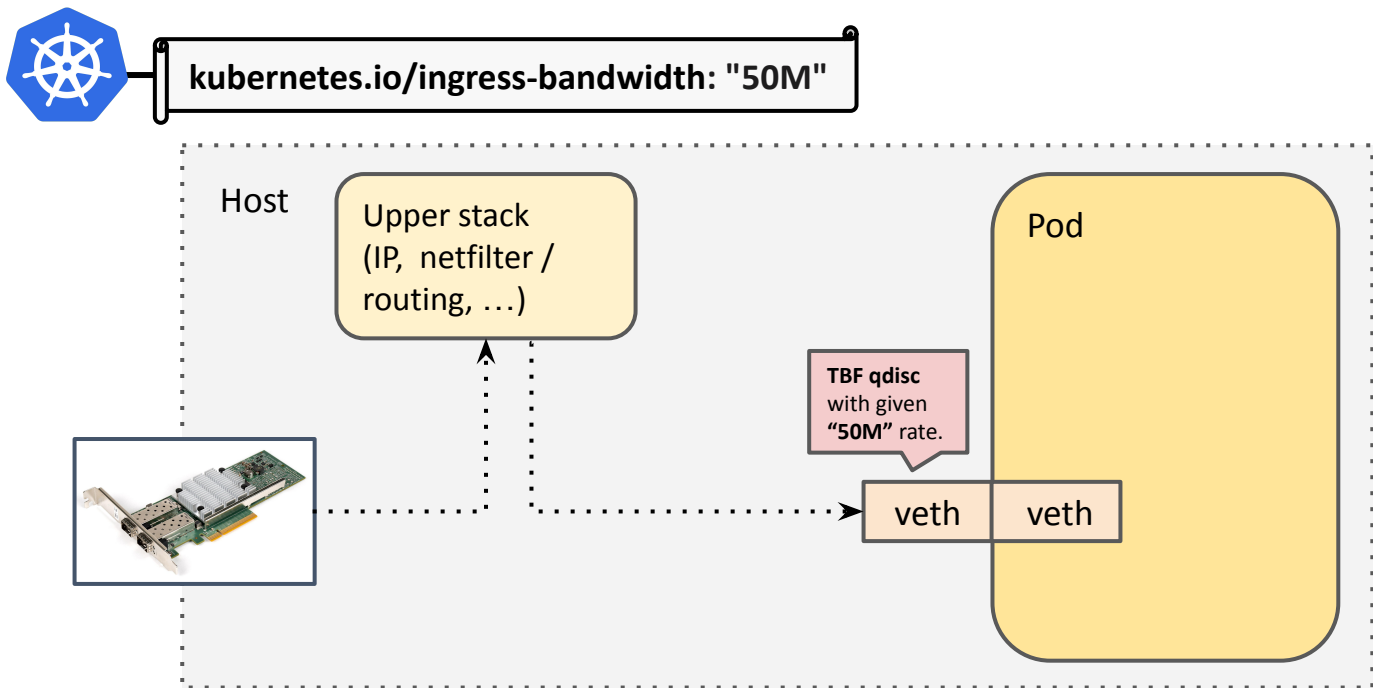
- bandwidth meta plugin not scalable for production use
- TBFs are attached to the Pod's veth devices



Problem Statement



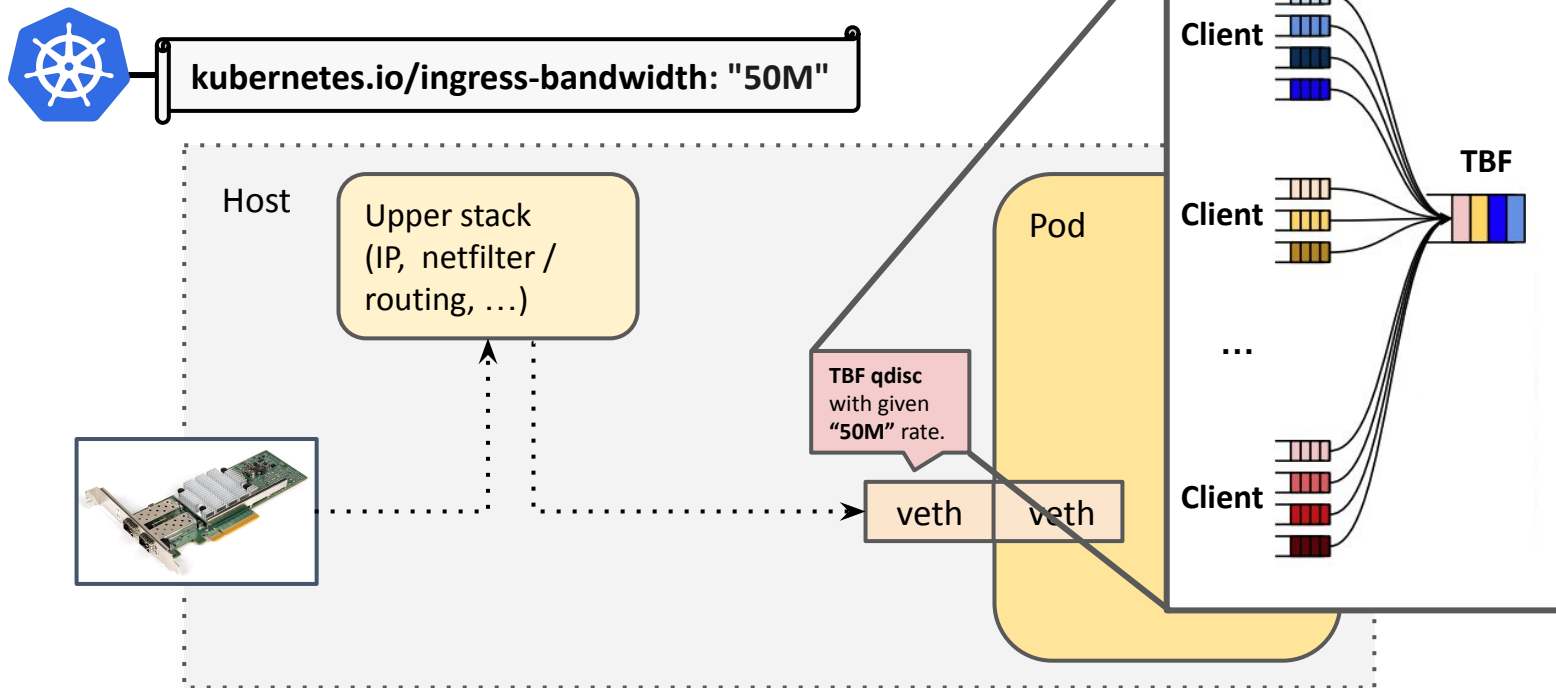
→ bandwidth meta plugin ingress example:



Problem Statement



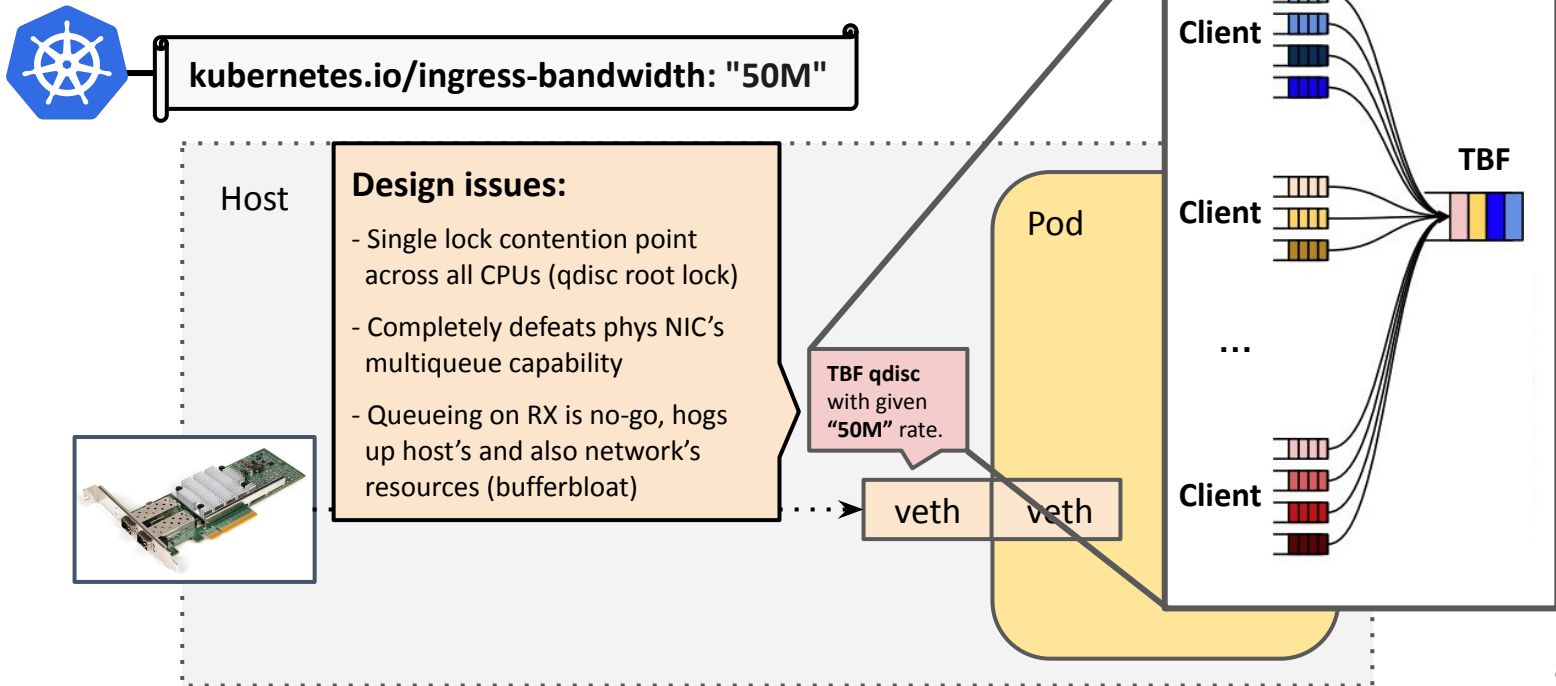
→ bandwidth meta plugin ingress example:



Problem Statement

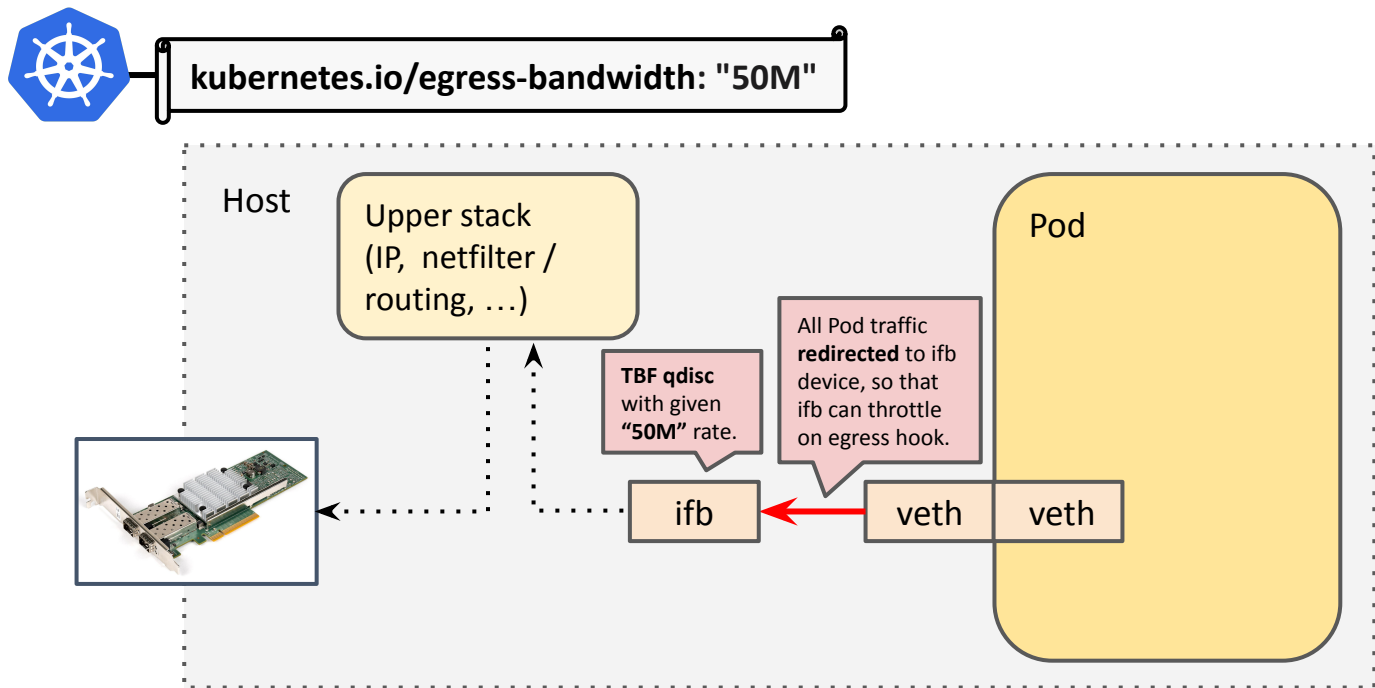


→ bandwidth meta plugin ingress example:



Problem Statement

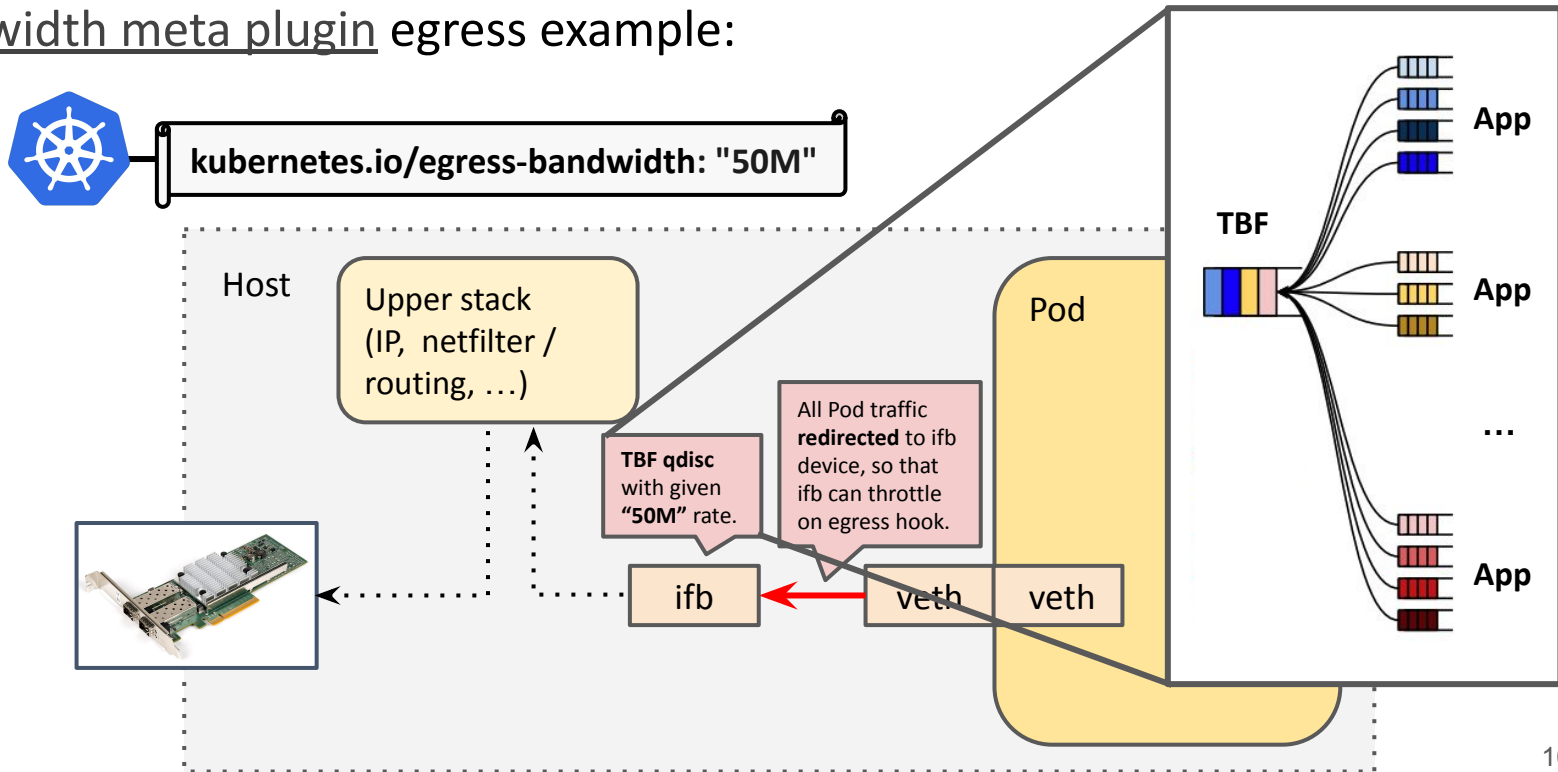
→ bandwidth meta plugin egress example:



Problem Statement



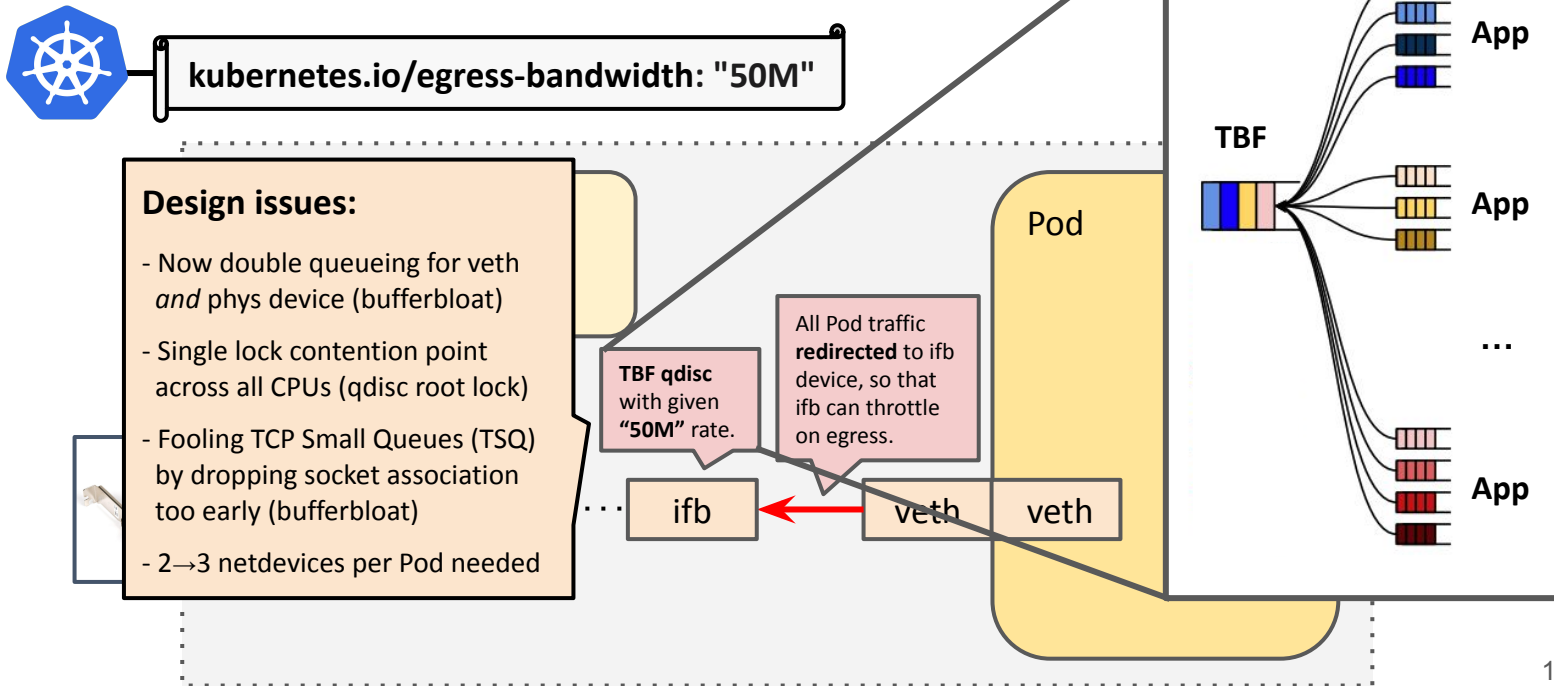
→ bandwidth meta plugin egress example:



Problem Statement



→ bandwidth meta plugin egress example:



Problem Statement



- bandwidth meta plugin aka “latency killer”
- tl;dr summary:



Problem Statement

... but AFAP makes bottleneck run at 100%

Queuing theory says this is fragile. E.g., for M/D/1:


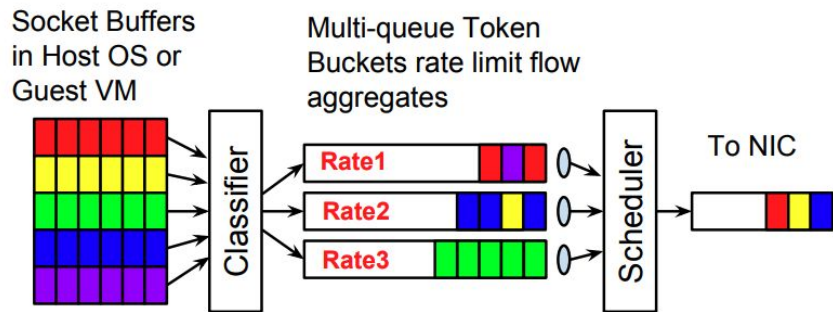


Packet wait-time in queue skyrockets when bottleneck link gets utilized close to 100%.

From Queues to EDT model

Core Idea to replace queues with two simple pieces:

- Earliest Departure Time (EDT) time stamp in every packet
- Timing-wheel scheduler which replaces the queue



How can the EDT model be applied to Kubernetes?



Programmable and **performant** in-kernel “virtual machine” that **safely** executes native code on certain events/hooks (aka “JavaScript for the kernel”).

How can the EDT model be applied to Kubernetes?

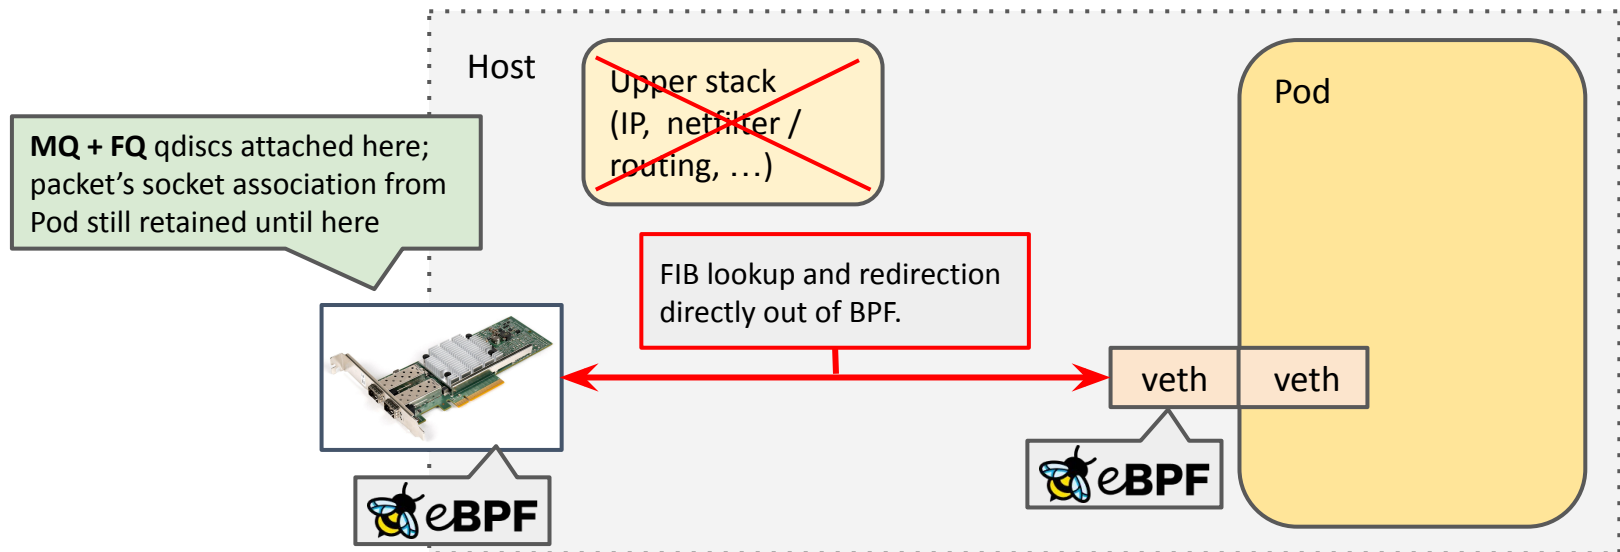


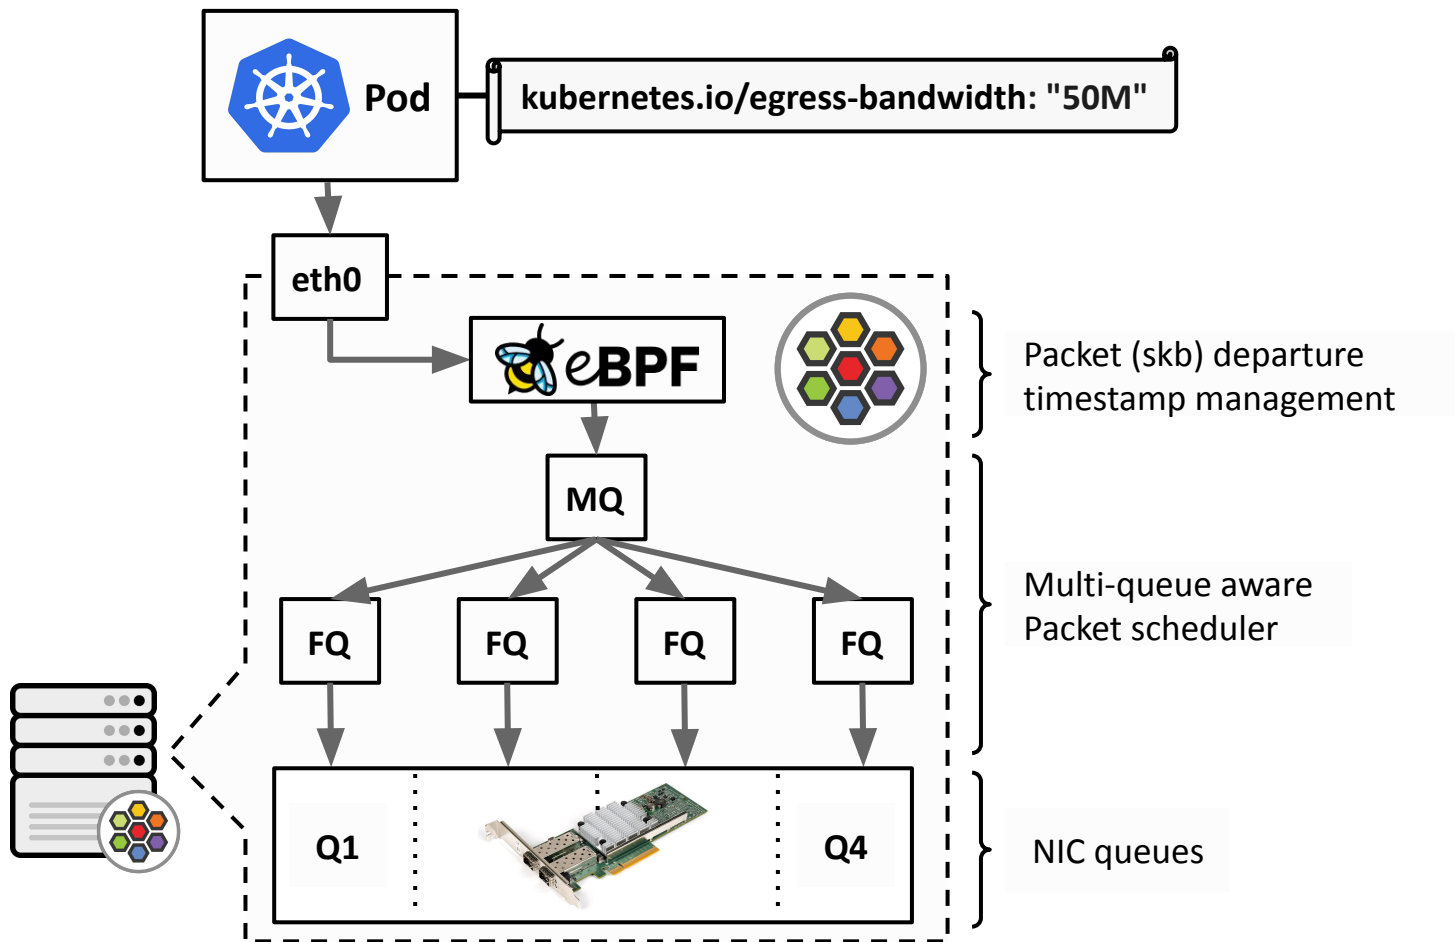
eBPF-based CNI / platform which provides Pod connectivity, service load-balancing, network policies, **bandwidth management**, transparent encryption and more.

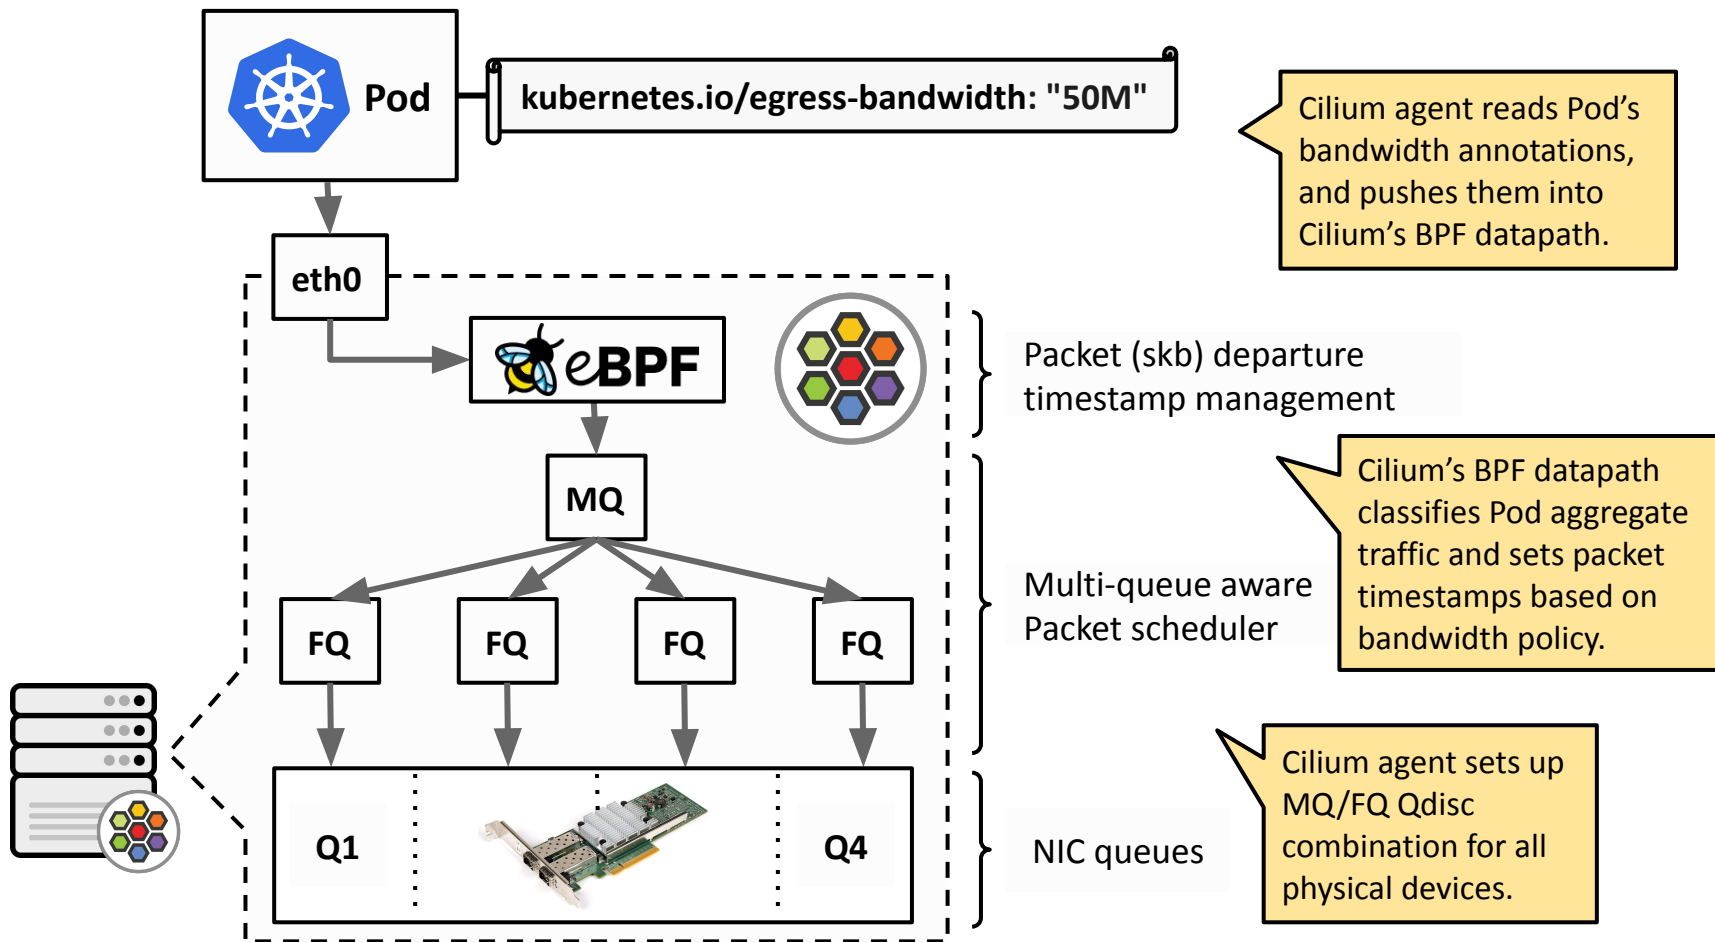


Cilium's Bandwidth Manager

- Implements lock-less, EDT-based Pod rate-limiting with eBPF
- Enforcement points on phys devices instead of veths to avoid bufferbloat and improve TCP TSQ feedback








Comparison of Cilium's EDT implementation vs TBF

Single flow latency for EDT and HTB/TBF model (lower is better)



Env: 256 concurrent request/response type flows (TCP_RR), 100M rate per flow




Comparison of Cilium's EDT implementation vs TBF

Total transaction rate between EDT and HTB/TBF model (higher is better)

■ EDT ■ TBF/HTB

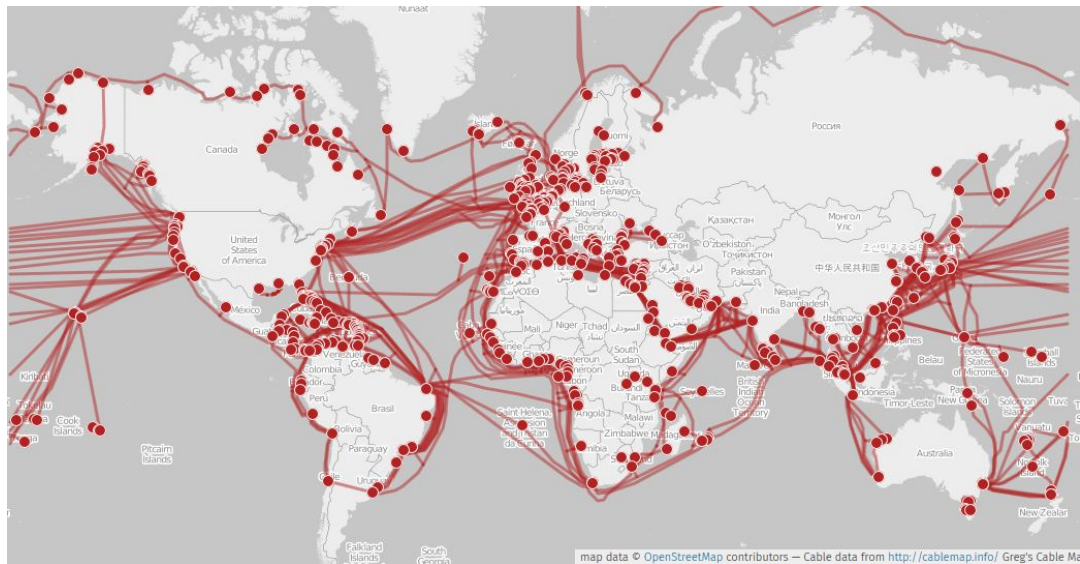
Env: 256 concurrent
request/response
type flows (TCP_RR),
100M rate per flow





Cilium's Bandwidth Manager: recap for now


- So far: Cilium's EDT approach allows for scalable bandwidth enforcement
- What about more broadly Internet-level bandwidth management?



What else does EDT model enable? Enter: BBR


TCP before BBR

Today's Internet is not moving data as well as it should. TCP sends data at lower bandwidth because the 1980s-era algorithm assumes that packet loss means network congestion.



TCP BBR

BBR models the network to send as fast as the available bandwidth and is 2700x faster than previous TCPs on a 10Gb, 100ms link with 1% loss. BBR powers google.com, youtube.com, and apps using Google Cloud Platform services.




Source: <https://cloud.google.com/blog/products/networking/tcp-bbr-congestion-control-comes-to-gcp-your-internet-just-got-faster>

What else does EDT model enable? Enter: BBR

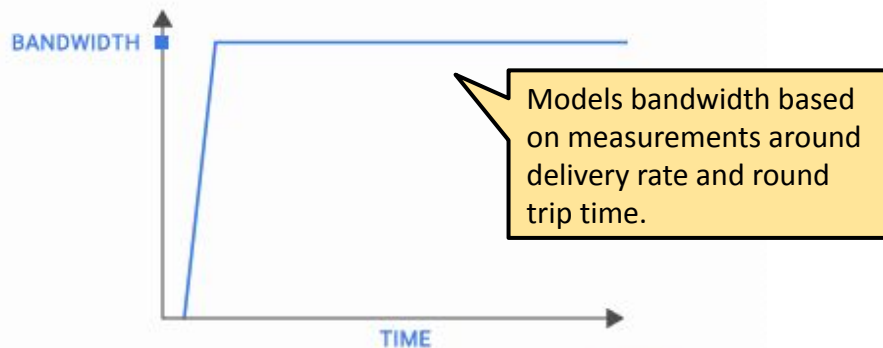
TCP before BBR

Today's Internet is not moving data as well as it should. TCP sends data at lower bandwidth because the 1980s-era algorithm assumes that packet loss means network congestion.



TCP BBR

BBR models the network to send as fast as the available bandwidth and is 2700x faster than previous TCPs on a 10Gb, 100ms link with 1% loss. BBR powers google.com, youtube.com, and apps using Google Cloud Platform services.





When is it useful to consider BBR?

- Kubernetes cluster exposing services to clients over the Internet
- ◆ Significant latency improvements for low-end last-mile networks
 - ◆ Significant throughput improvements for high-speed long-haul links

	CUBIC (default)	BBR (v1)
Model parameters to the state machine	N/A	Throughput, RTT
Loss	Reduce cwnd by 30% on window with any loss	N/A
ECN	RFC3168 (classic ECN)	N/A
Startup	Slow-start until RTT rises (Hystart) or any loss	Slow-start until throughput plateaus

Example: New York (packet.net) -> Zurich



```
darkstar@linux:~/trees/bpf-... x root@zh-lab-node-1: ~ x darkstar@linux:~/trees/bpf x root@ny-c3-small-x86-01: ~ x root@
root@zh-lab-node-1:~# while [ 1 ]; do iperf3 -c 147.75.66.15 -t 55 -R -i 5 -O 5 ; done
Connecting to host 147.75.66.15, port 5201
Reverse mode, remote host 147.75.66.15 is sending
[ 5] local 192.168.178.91 port 52148 connected to 147.75.66.15 port 5201
[ ID] Interval            Transfer        Bitrate
[ 5]  0.00-5.00      sec  78.4 MBytes    132 Mbits/sec
[ 5]  0.00-5.00      sec  98.2 MBytes    165 Mbits/sec
[ 5]  5.00-10.00     sec  99.1 MBytes    166 Mbits/sec
[ 5] 10.00-15.00     sec  113 MBytes     189 Mbits/sec
[ 5] 15.00-20.00     sec  159 MBytes     267 Mbits/sec
[ 5] 20.00-25.00     sec  257 MBytes     431 Mbits/sec
[ 5] 25.00-30.00     sec  153 MBytes     256 Mbits/sec
[ 5] 30.00-35.00     sec  146 MBytes     245 Mbits/sec
[ 5] 35.00-40.00     sec  148 MBytes     248 Mbits/sec
[ 5] 40.00-45.00     sec  157 MBytes     264 Mbits/sec
[ 5] 45.00-50.00     sec  193 MBytes     323 Mbits/sec
[ 5] 50.00-55.00     sec  272 MBytes     457 Mbits/sec
- - - - -
[ ID] Interval            Transfer        Bitrate      Retr
[ 5]  0.00-55.10     sec  1.76 GBytes    274 Mbits/sec  1501
[ 5]  0.00-55.00     sec  1.75 GBytes    274 Mbits/sec
```

(omitted)

Bandwidth probing, overreaction to loss!
(Sawtooth pattern nicely visible)

Default, server runs:

- TCP CUBIC
- fq_codel Qdisc

sender
receiver

Example: New York (packet.net) -> Zurich



```
darkstar@linux:~/trees/bpf-... x root@zh-lab-node-1: ~ x darkstar@linux:~/trees/bpf x root@ny-c3-small-x86-01: ~ x root@
root@zh-lab-node-1:~# while [ 1 ]; do iperf3 -c 147.75.66.15 -t 55 -R -i 5 -O 5 ; done
Connecting to host 147.75.66.15, port 5201
Reverse mode, remote host 147.75.66.15 is sending
[ 5] local 192.168.178.91 port 52254 connected to 147.75.66.15 port 5201
[ ID] Interval          Transfer      Bitrate
[ 5]  0.00-5.00      sec    152 MBytes   254 Mbits/sec
[ 5]  0.00-5.00      sec    258 MBytes   433 Mbits/sec
[ 5]  5.00-10.00     sec    240 MBytes   403 Mbits/sec
[ 5] 10.00-15.00     sec    255 MBytes   427 Mbits/sec
[ 5] 15.00-20.00     sec    247 MBytes   414 Mbits/sec
[ 5] 20.00-25.00     sec    255 MBytes   428 Mbits/sec
[ 5] 25.00-30.00     sec    255 MBytes   428 Mbits/sec
[ 5] 30.00-35.00     sec    238 MBytes   400 Mbits/sec
[ 5] 35.00-40.00     sec    255 MBytes   428 Mbits/sec
[ 5] 40.00-45.00     sec    239 MBytes   401 Mbits/sec
[ 5] 45.00-50.00     sec    253 MBytes   425 Mbits/sec
[ 5] 50.00-55.00     sec    242 MBytes   407 Mbits/sec
-- -- -- -- --
[ ID] Interval          Transfer      Bitrate      Retr
[ 5]  0.00-55.10     sec    2.68 GBytes   418 Mbits/sec 58812
[ 5]  0.00-55.00     sec    2.67 GBytes   418 Mbits/sec
```

No overreaction, stable.

Updated, server runs:

- TCP BBR
- FQ Qdisc (for EDT)

(results also reproduce with netperf)

sender
receiver

Example: New York (packet.net) -> Zurich



```
darkstar@linux:~/trees/bpf-... x root@zh-lab-node-1: ~ x darkstar@linux:~/trees/bpf x root@ny-c3-small-x86-01: ~ x root@
```

```
root@zh-lab-node-1:~# while [ 1 ]; do iperf3 -c 147.75.66.15 -t 55 -R -i 5 -O 5 ; done
```

```
Connecting to host 147.75.66.15, port 5201
```

```
Reverse mode, remote host 147.75.66.15 is sending
```

```
[ 5] local 192.168.178.91 port 52254 connected to 147.75.66.15 port 5201
```

[ID]	Interval		Transfer	Bitrate	
[5]	0.00-5.00	sec	152 MBytes	254 Mbits/sec	(omitted)
[5]	0.00-5.00	sec	258 MBytes	433 Mbits/sec	
[5]	5.00-10.00	sec	240 MBytes	403 Mbits/sec	
[5]	10.00-15.00	sec	255 MBytes	427 Mbits/sec	
[5]	15.00-20.00	sec	247 MBytes	414 Mbits/sec	
[5]	20.00-25.00	sec	255 MBytes	428 Mbits/sec	
[5]	25.00-30.00	sec	255 MBytes	428 Mbits/sec	
[5]	30.00-35.00	sec	238 MBytes	400 Mbits/sec	
[5]	35.00-40.00	sec	255 MBytes	428 Mbits/sec	
[5]	40.00-45.00	sec	239 MBytes	401 Mbits/sec	
[5]	45.00-50.00	sec	253 MBytes	425 Mbits/sec	
[5]	50.00-55.00	sec	242 MBytes	407 Mbits/sec	

[ID]	Interval		Transfer	Bitrate	Retr
[5]	0.00-55.10	sec	2.68 GBytes	418 Mbits/sec	58812
[5]	0.00-55.00	sec	2.67 GBytes	418 Mbits/sec	

No overreaction, stable.

Updated, server runs:

- TCP BBR
- FQ Qdisc (for EDT)

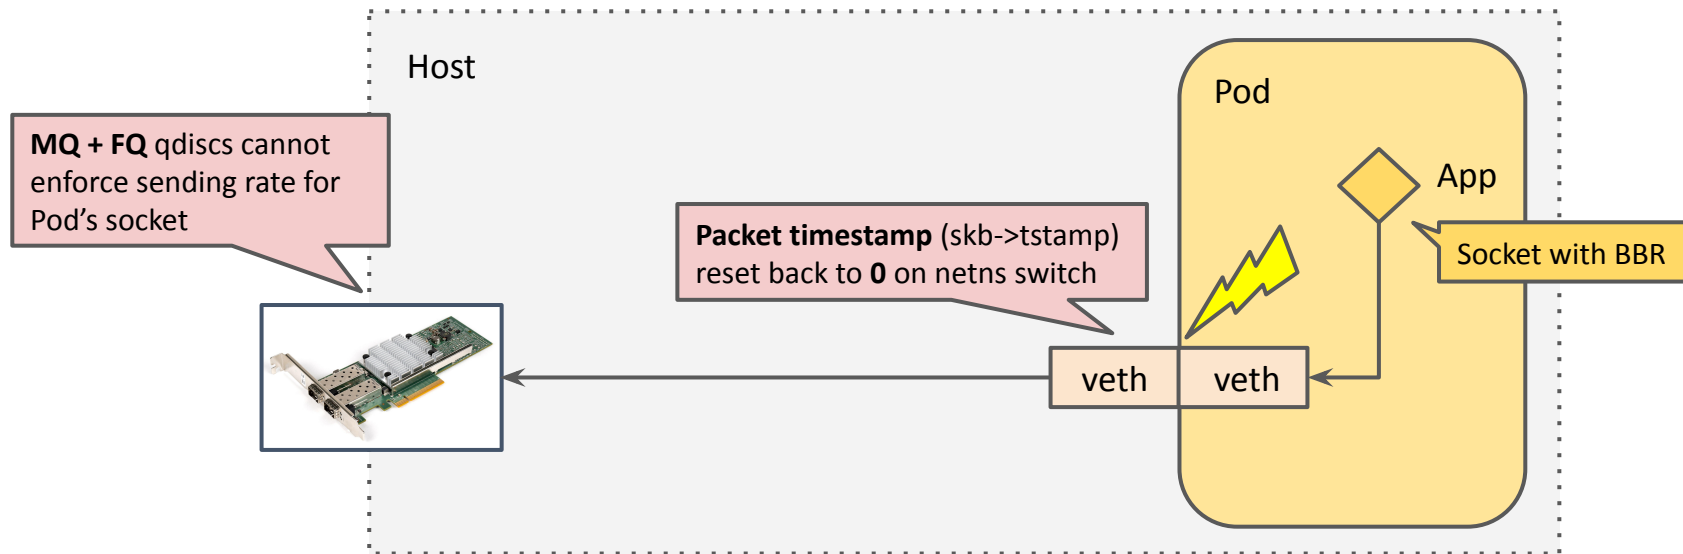
CUBIC → **BBR** bumps:

274 → **418 Mbit/s**




Can BBR be used for Kubernetes Pods?

- BBR works in conjunction with FQ and sets packet delivery timestamps
- Kernel clears timestamp for packets leaving Pods (== netns)
 - ◆ Usage of BBR for Pods not possible/broken in general today



Rationale on today's timestamp reset ([see our LPC talk](#))



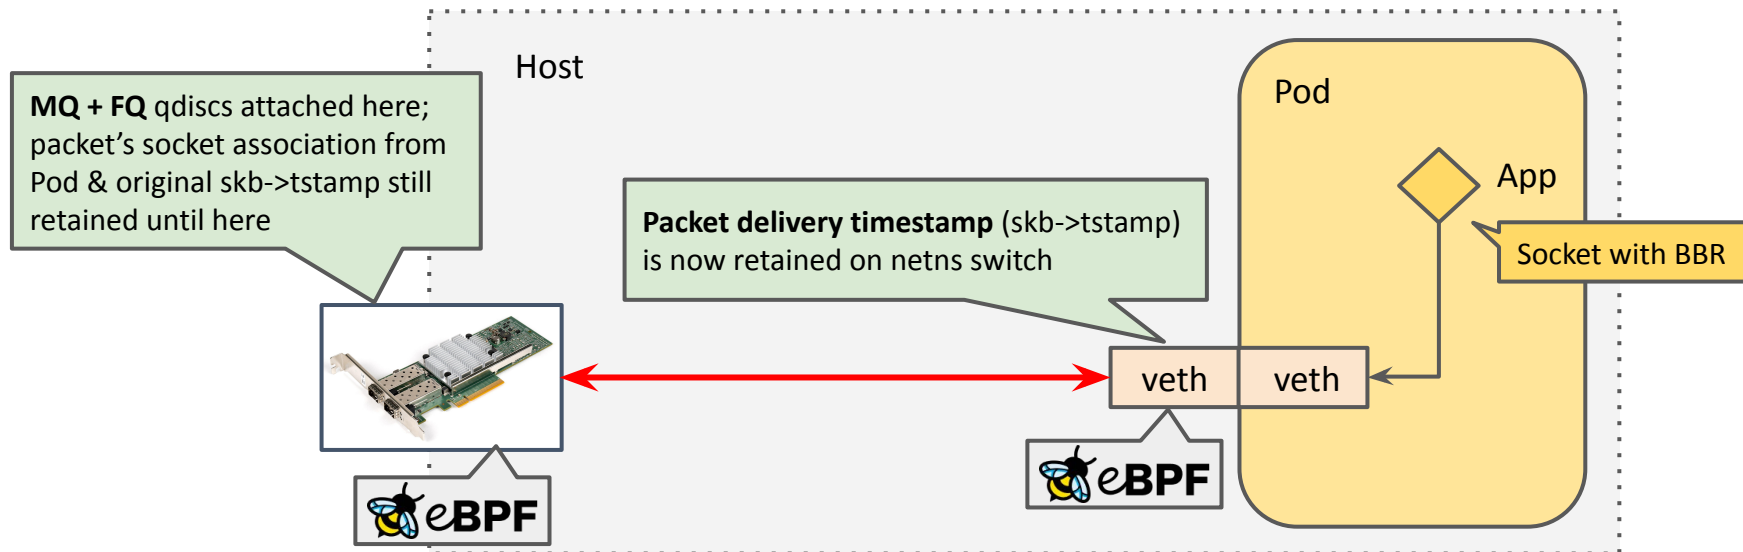
Kernel uses different clock bases for `skb->timestamp`:

- Ingress is `CLOCK_TAI`, egress is `CLOCK_MONOTONIC` (as is FQ)
- Forwarding from RX to TX would cause drop in FQ due to overreaching FQ's drop horizon given different clock's offsets
- No means to figure out clock base from `skb->timestamp`, hence reset



Cilium's Bandwidth Manager: BBR

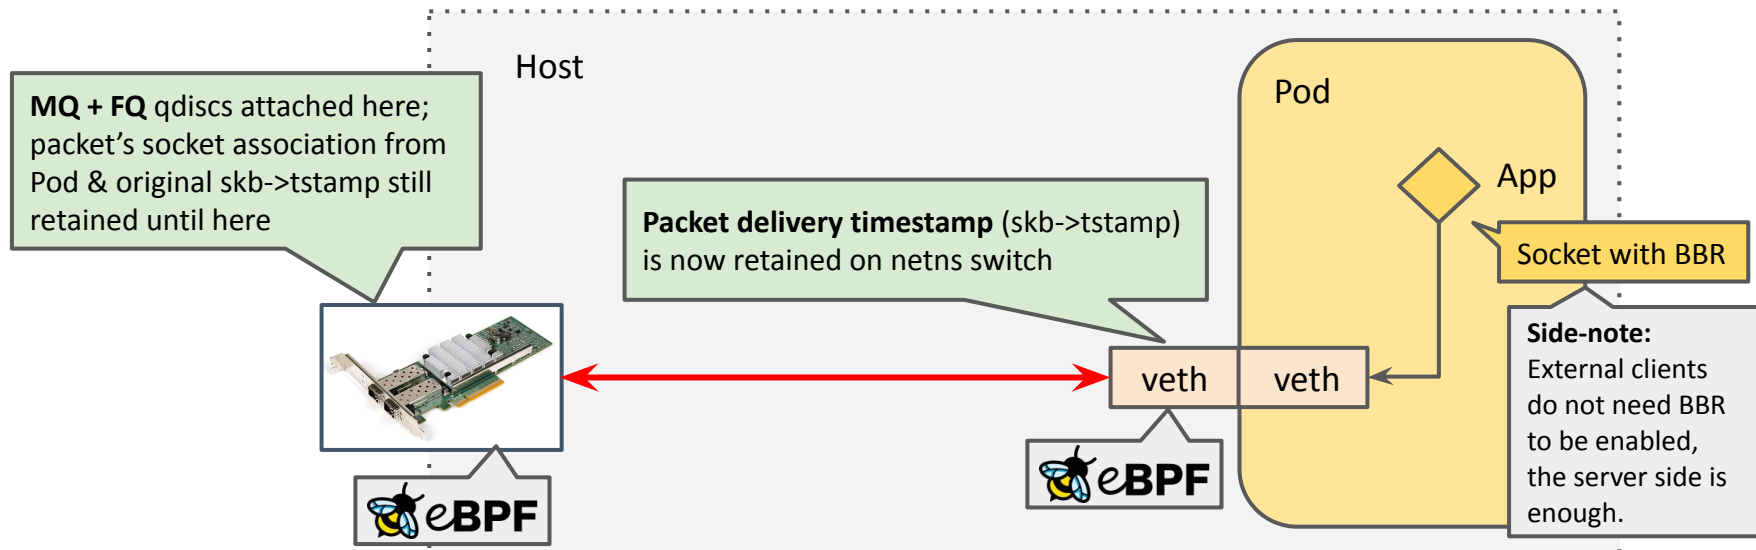
- We helped fixing networking stack in Linux v5.18+ to retain timestamps
- Bandwidth Manager plumbs the appropriate underlying infrastructure
 - ◆ Receives new knob for switching whole cluster over to BBR by default





Cilium's Bandwidth Manager: BBR

- We helped fixing networking stack in Linux v5.18+ to retain timestamps
- Bandwidth Manager plumbs the appropriate underlying infrastructure
 - ◆ Receives new knob for switching whole cluster over to BBR by default




ISOVALENT

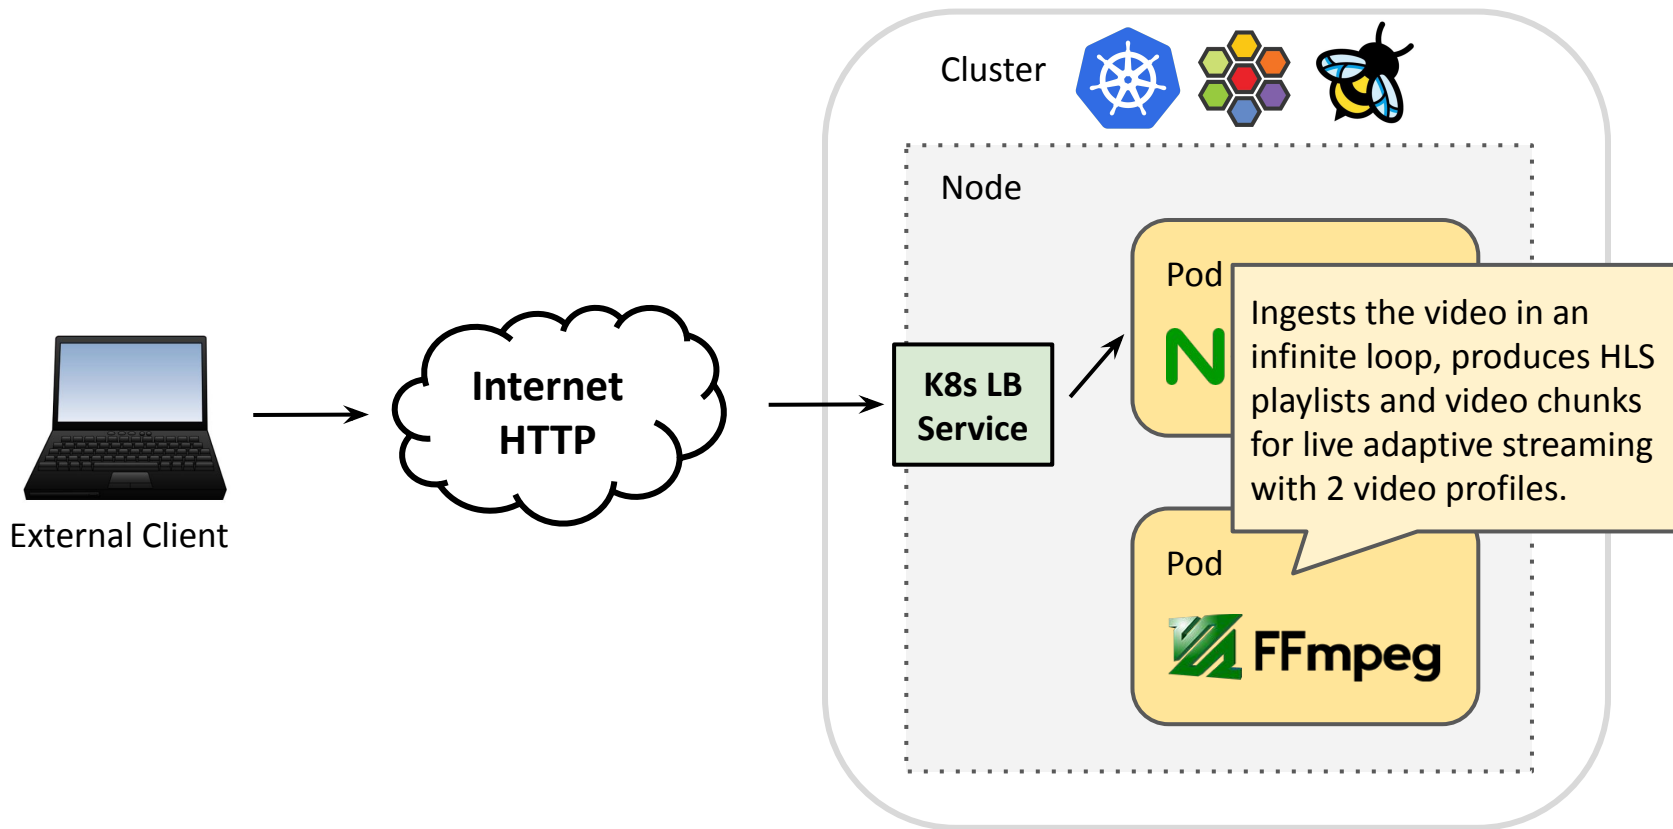
Demo: BBR for Pods

(K8s/Cilium-backed video streaming service: CUBIC versus BBR)

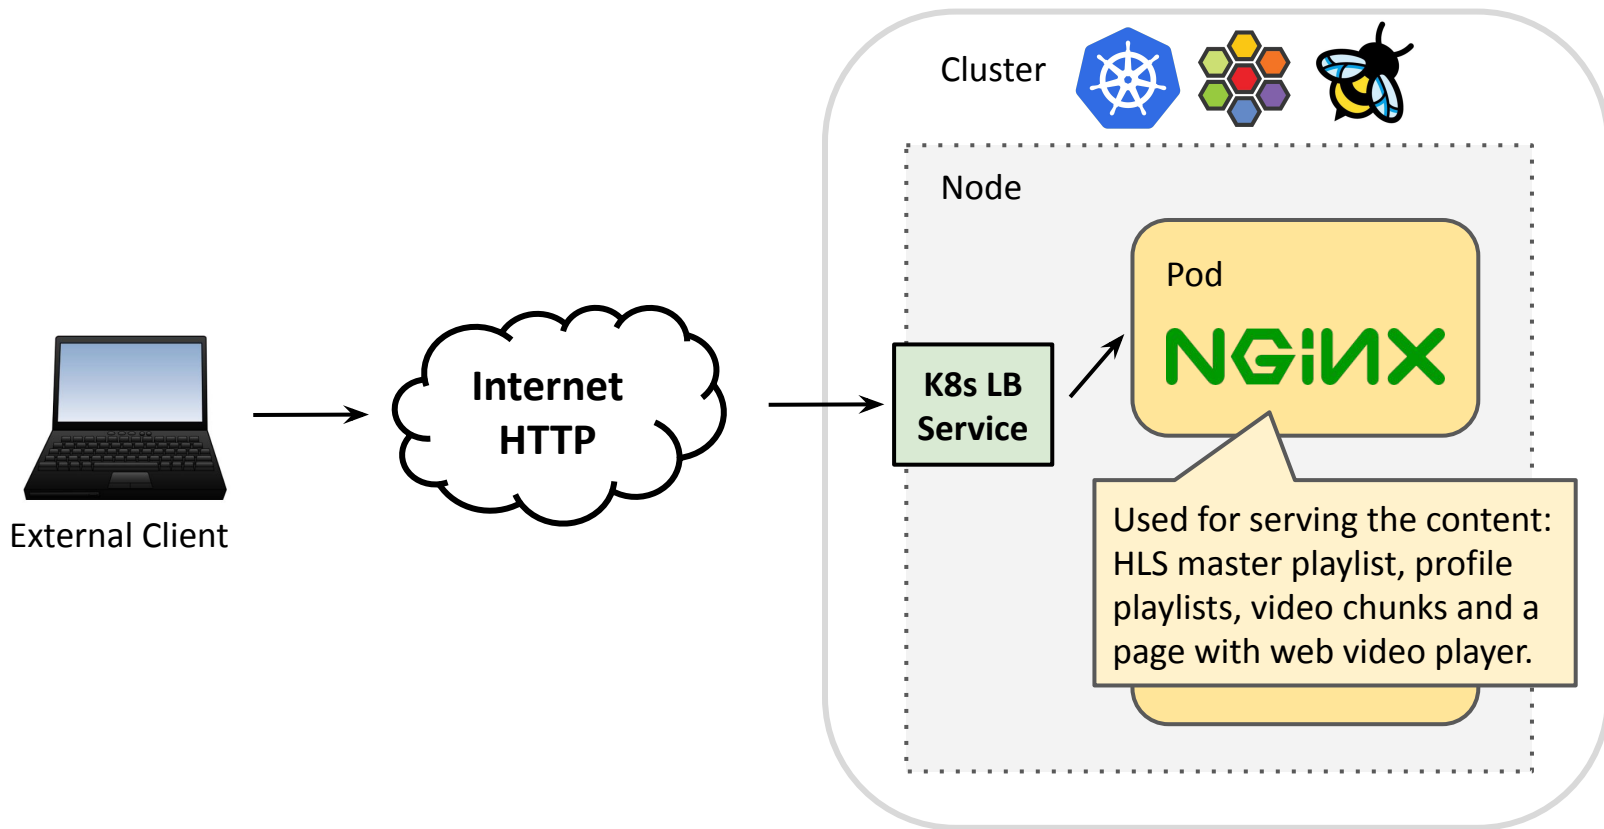
Video Streaming Service Demo Setup: CUBIC vs BBR



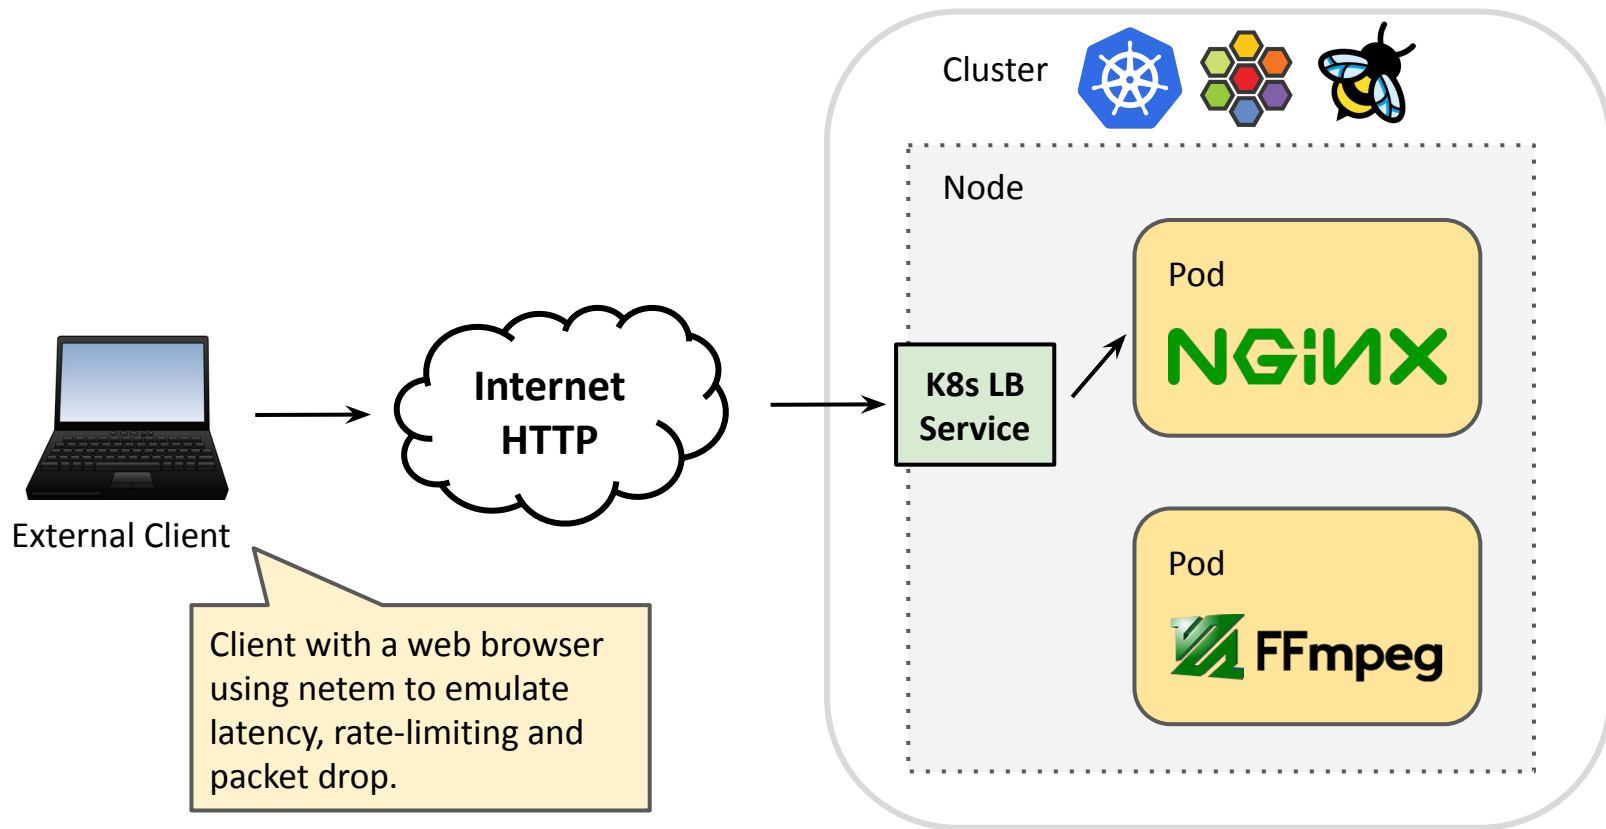
Video Streaming Service Demo Setup: CUBIC vs BBR



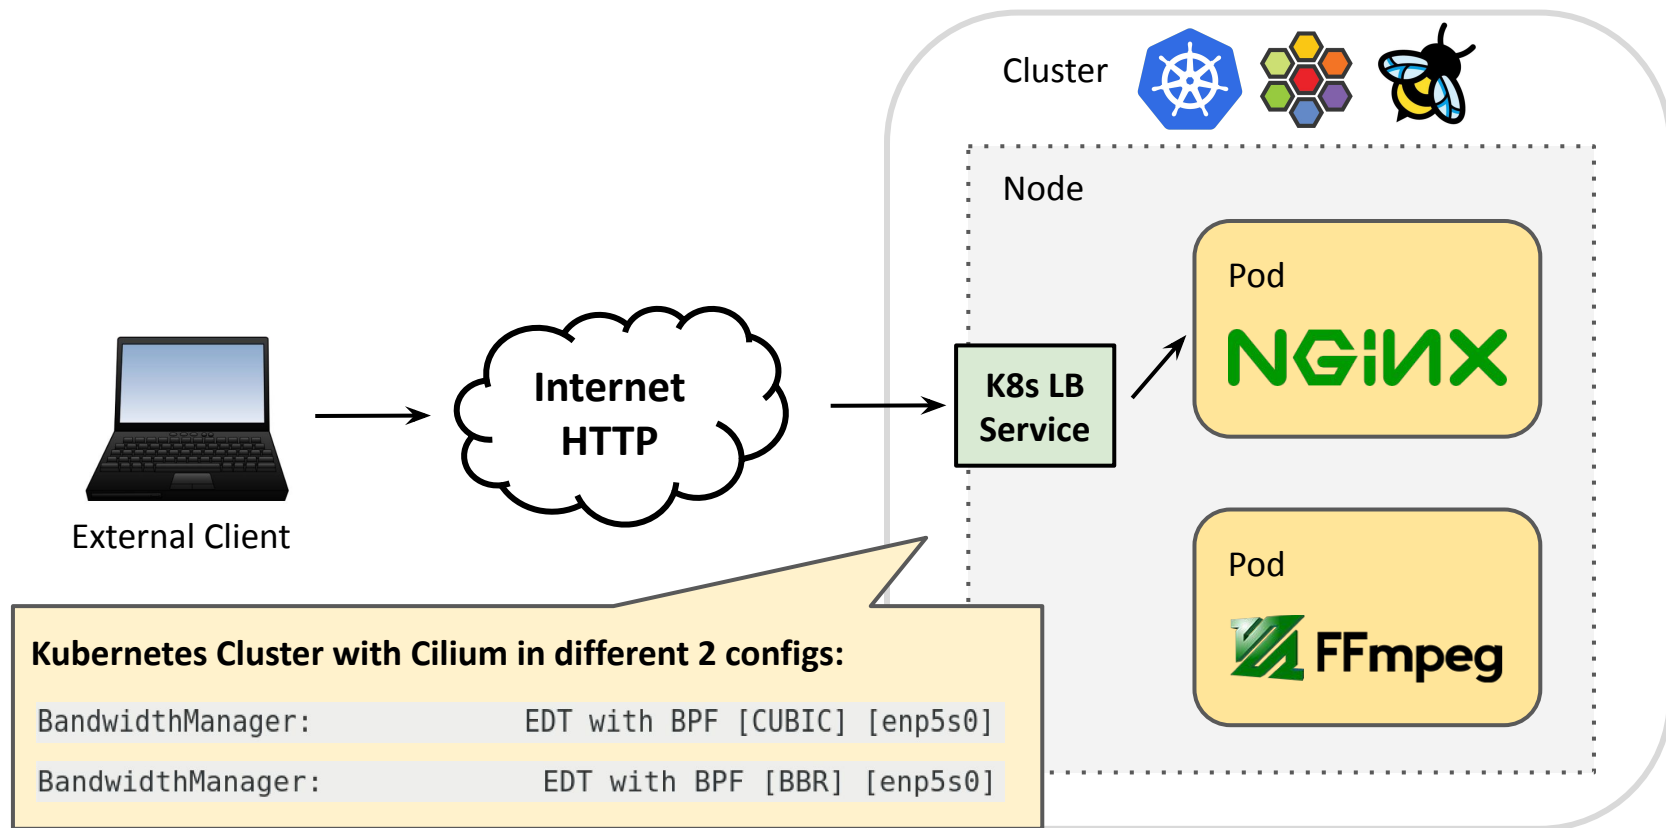
Video Streaming Service Demo Setup: CUBIC vs BBR



Video Streaming Service Demo Setup: CUBIC vs BBR



Video Streaming Service Demo Setup: CUBIC vs BBR



ISOVALENT

A photograph of a theater interior. In the foreground, rows of plush red seats are visible, receding into the distance. The stage is at the far end, covered by a large, dark, vertically-pleated curtain. The lighting is dim, creating a sense of anticipation. The text "Now to the stream ..." is overlaid in the center of the image.

Now to the stream ...



What needs to be considered with use of BBR?

- BBR has potential unfairness issues towards CUBIC when env uses both
- BBR will trigger a higher TCP retransmission rate (more aggressive probing)
- BBRv2 in the works to overcome them





Revisiting earlier Problem Statement

- Kubernetes bandwidth enforcement does not need to be in a poor state
- Native implementation via Cilium's Bandwidth Manager (GA since v1.12)
 - ◆ Efficient, eBPF-based bandwidth enforcement via EDT model
 - ◆ First CNI to support BBR (& socket pacing) for Pods
 - ◆ Side-note: Realizing such architecture only possible with eBPF



Getting Started Guide for Bandwidth Manager:

```
helm upgrade cilium ./cilium \  
  --namespace kube-system \  
  --reuse-values \  
  --set bandwidthManager.enabled=true \  
  --set bandwidthManager.bbr=true  
kubectl -n kube-system rollout restart ds/cilium
```

(needs Linux kernel v5.1+)

(needs Linux kernel v5.18+)

Acknowledgements

- Van Jacobson
- Eric Dumazet
- Vytautas Valancius
- Stanislav Fomichev
- Martin Lau
- John Fastabend
- Cilium, BPF & netdev kernel community

ISOVALENT

Thank you! Questions?

github.com/cilium/cilium

cilium.io

ebpf.io

Isovalent: booth S21


Cilium: kiosk 12 (@ project pavilion)



Appendix: Latency & TPS for 64/128/256 flows

Comparison of Cilium's EDT implementation vs TBF

Single flow latency for EDT and HTB/TBF model (lower is better)




Env: 64 concurrent request/response type flows (TCP_RR), 100M rate per flow



Comparison of Cilium's EDT implementation vs TBF

Single flow latency for EDT and HTB/TBF model (lower is better)




Env: 128 concurrent request/response type flows (TCP_RR), 100M rate per flow



Comparison of Cilium's EDT implementation vs TBF

Single flow latency for EDT and HTB/TBF model (lower is better)

■ EDT ■ TBF/HTB



Env: 256 concurrent request/response type flows (TCP_RR), 100M rate per flow




Comparison of Cilium's EDT implementation vs TBF

Total transaction rate between EDT and HTB/TBF model (higher is better)

■ EDT ■ TBF/HTB


Env: 64 concurrent
request/response
type flows (TCP_RR),
100M rate per flow



Comparison of Cilium's EDT implementation vs TBF

Total transaction rate between EDT and HTB/TBF model (higher is better)

■ EDT ■ TBF/HTB



Comparison of Cilium's EDT implementation vs TBF

Total transaction rate between EDT and HTB/TBF model (higher is better)

■ EDT ■ TBF/HTB

