

Kubernetes Everywhere

Lessons Learned From Going Multi-Cloud

20 May 2022



Niko Smeds

Cloud Platform at Grafana



Kubernetes Everywhere

Agenda

1. Why go multi-cloud?
2. Project overview
3. Five lessons
4. What went well
5. Q&A




Why?

Why might an organization consider opting for multi-cloud?

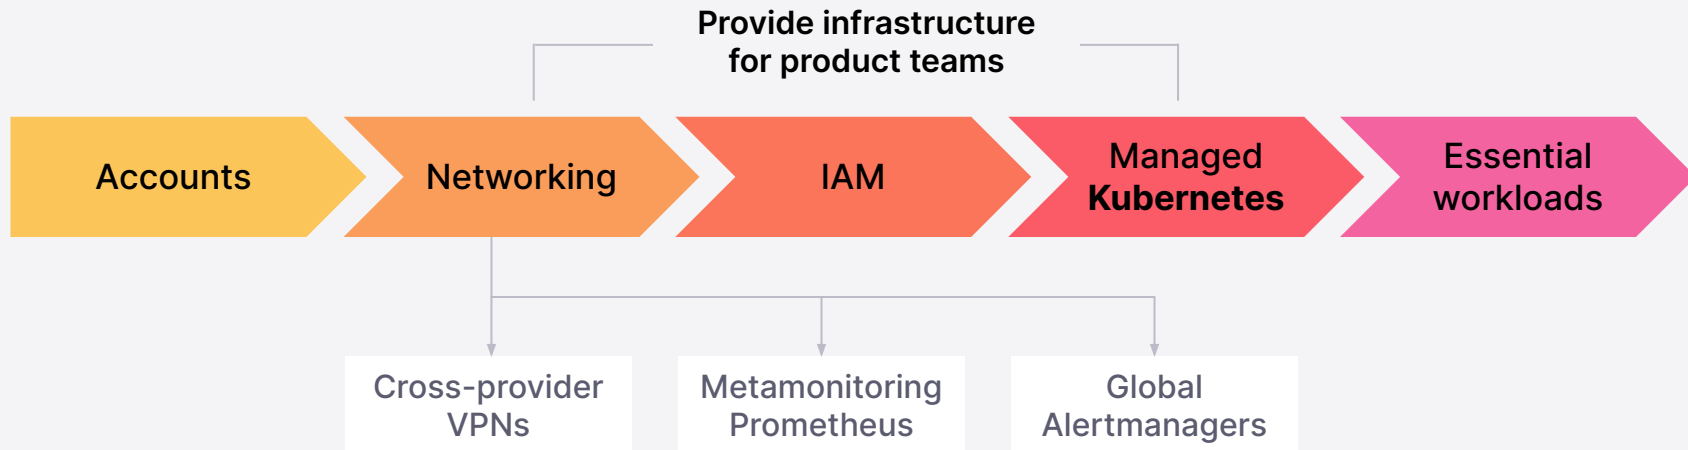
- Increase **available regions**
- Reduce **vendor lock-in**
- Customer **preference**
 - Latency
 - Data sovereignty
 - Spend commit



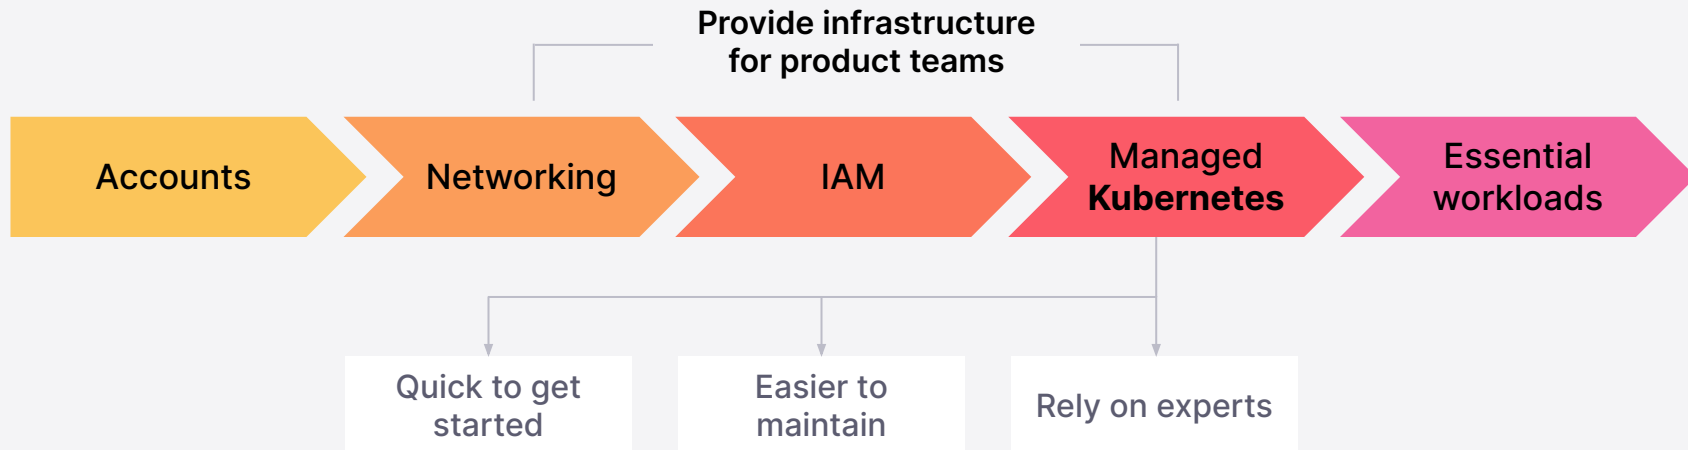
Project: Grafana Cloud Platform on AWS



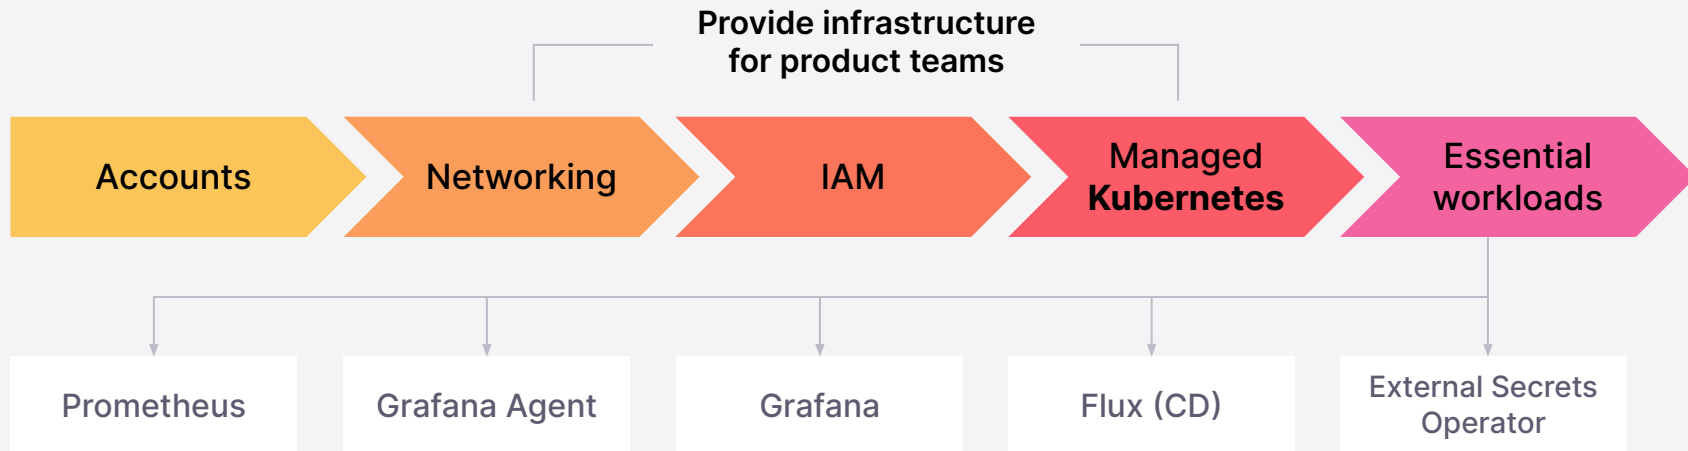
Project: Grafana Cloud Platform on AWS



Project: Grafana Cloud Platform on AWS



Project: Grafana Cloud Platform on AWS



Getting Started

In the beginning...

- Set **clear requirements**
 - What workloads to include/exclude?
 - Do we avoid inter-provider dependencies?
 - Expected scale?
- **Research** Amazon Elastic Kubernetes (EKS)
 - Documentation
 - Conference videos
 - Blog posts



READY,

SET,

... not quite



Virtual Private Cloud (VPC)

GCP

Global VPCs

Supports subnets from multiple regions.

We have use-cases for this and run multiple K8s clusters across multiple regions in shared VPCs.

vs

AWS

Regional VPCs

Supports subnets from a single region.

We're left to decide:

- VPC per region, or
- VPC per cluster



VPC CIDRs

GCP

Mix private ranges

Alias IP ranges supported from a mix of RFC 1918 ranges.

- **10.0.0.0/8** subnets for primary components (nodes, pods, services)
- **172.16.0.0/12** subnets for managed components (GKE control plane)



AWS

Single private range

Secondary CIDR block restricted to a single RFC 1918 range.

- **10.0.0.0/8** subnets for primary components (nodes, pods, services)
- **10.0.0.0/8** subnets for managed components (EKS control plane)

RFC 1918 ranges

10.0.0.0/8
172.16.0.0/12
192.168.0.0/16



Subnets

GCP

Supports /29 → /8

Pods are deployed in **/14** subnets.

vs

AWS

Supports /28 → /16

Pods are deployed in **/16** subnets.

Required refactoring our **IP reservation plans**.
We avoid private range overlap to support
inter-cluster peering.



Comparing other services



Load Balancers

Mix of global and non-global
load balancers.



Comparing other services



Load Balancers

Mix of global and non-global load balancers.



Volumes

Different tiers, performance, and pricing.



Comparing other services



Load Balancers

Mix of global and non-global load balancers.



Volumes

Different tiers, performance, and pricing.



Object Storage

Variance in read/write rate-limits.



1

Provider services
are similar, but not
the same



Tutorial hell



~~**Tutorial hell**~~

Documentation hell





*I hear and I forget.
I see and I remember.
I do and I understand.*

Xunzi
3rd Century BCE



Getting Started (again...)

Start building with the knowledge you've gathered.

- Be ready to **tear it down** and **start over**
- Use **infrastructure as code** and **version control** from the get-go
 - Peer review and suggestions
 - Share progress
 - Documents history
- You **cannot plan** for the **unknown**
 - You'll face limitations and bugs you didn't expect, and that's okay



2

Just get started



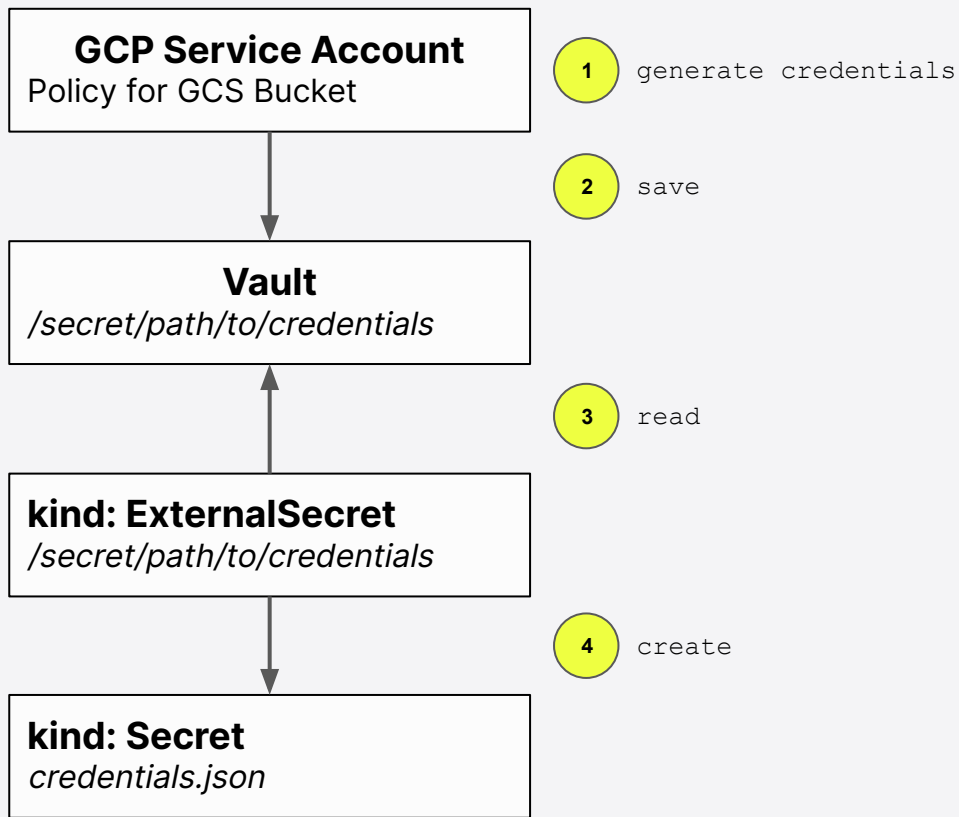
Expanding to a new cloud provider?

ctrl + c

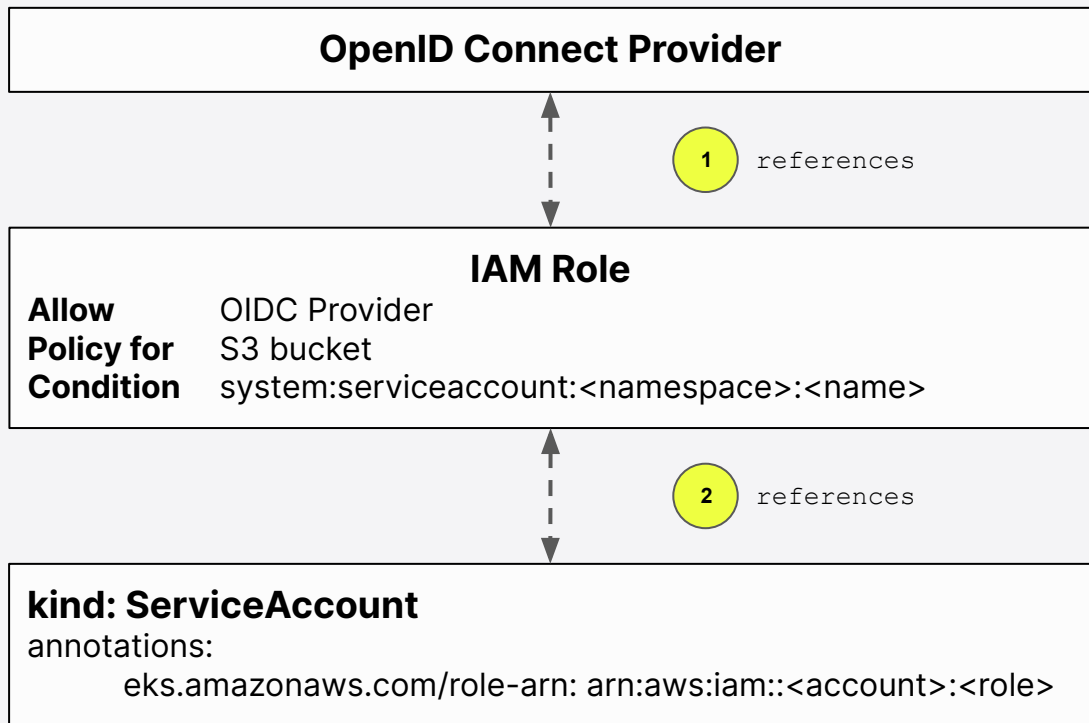
ctrl + v



Example: Application credentials (GCP)



Example: Application credentials (AWS)



3

Don't blindly do
what was done
before



It's about time



AWS **EKS cluster** is up and running.



Add-ons are installed:

- **Cluster Autoscaler**
- **Load Balancer Controller**
- **AWS VPC CNI**



The **essential workloads** are deployed.



READY,

SET,

... uh oh



Issue: Poor disk performance

Turns out we were using the **default EKS Storage Classes**.

- Previous generation volumes
- Slowest IOPS option
 - Min: **100 IOPS**
 - Max: 5000 IOPS

Resolved by:

- Install **Amazon EBS CSI driver**
- Upgrade Storage Classes



Issue: Docker Hub rate-limits

Pods failing to pull images due to Docker rate-limiting.

- Load tests triggered **pod scaling**
 - e.g. 10s of pods to 100s of pods
- AWS NAT Gateway uses a **single IP address** for egress
- GCP provides an **image cache** by default

Resolved by:

- Set up Docker registry mirror as **pull through cache**
- Update **node pools** to utilize the internal mirror



4

Load test your
applications



“

The cloud is infinite.



Niko

Early career days



“

The cloud is infinite.

MYTH



Niko

Early career days



Quotas

New accounts start with (small) default service limits.

- Regions supported by **physical data centers**
- We're **blind to capacity** of the regions we rely on
- Quotas increased via **support tickets**
 - Usually resolved within minutes
 - Sometimes days
 - Sometimes... not at all



5

Know your quotas



Recap

**Similar,
but not the
same**

**Just get
started**

**Don't blindly
do what was
done before**

**Load test
your
applications**

**Know your
quotas**



What went well?

- **Kubernetes**

- Most **Deployments**, **StatefulSets** and other workloads “just worked”
- Tweaks required for
 - **Ingress**
 - **Services** of type **LoadBalancer**
 - Applications/configs which referenced object storage buckets

○





Gracias

